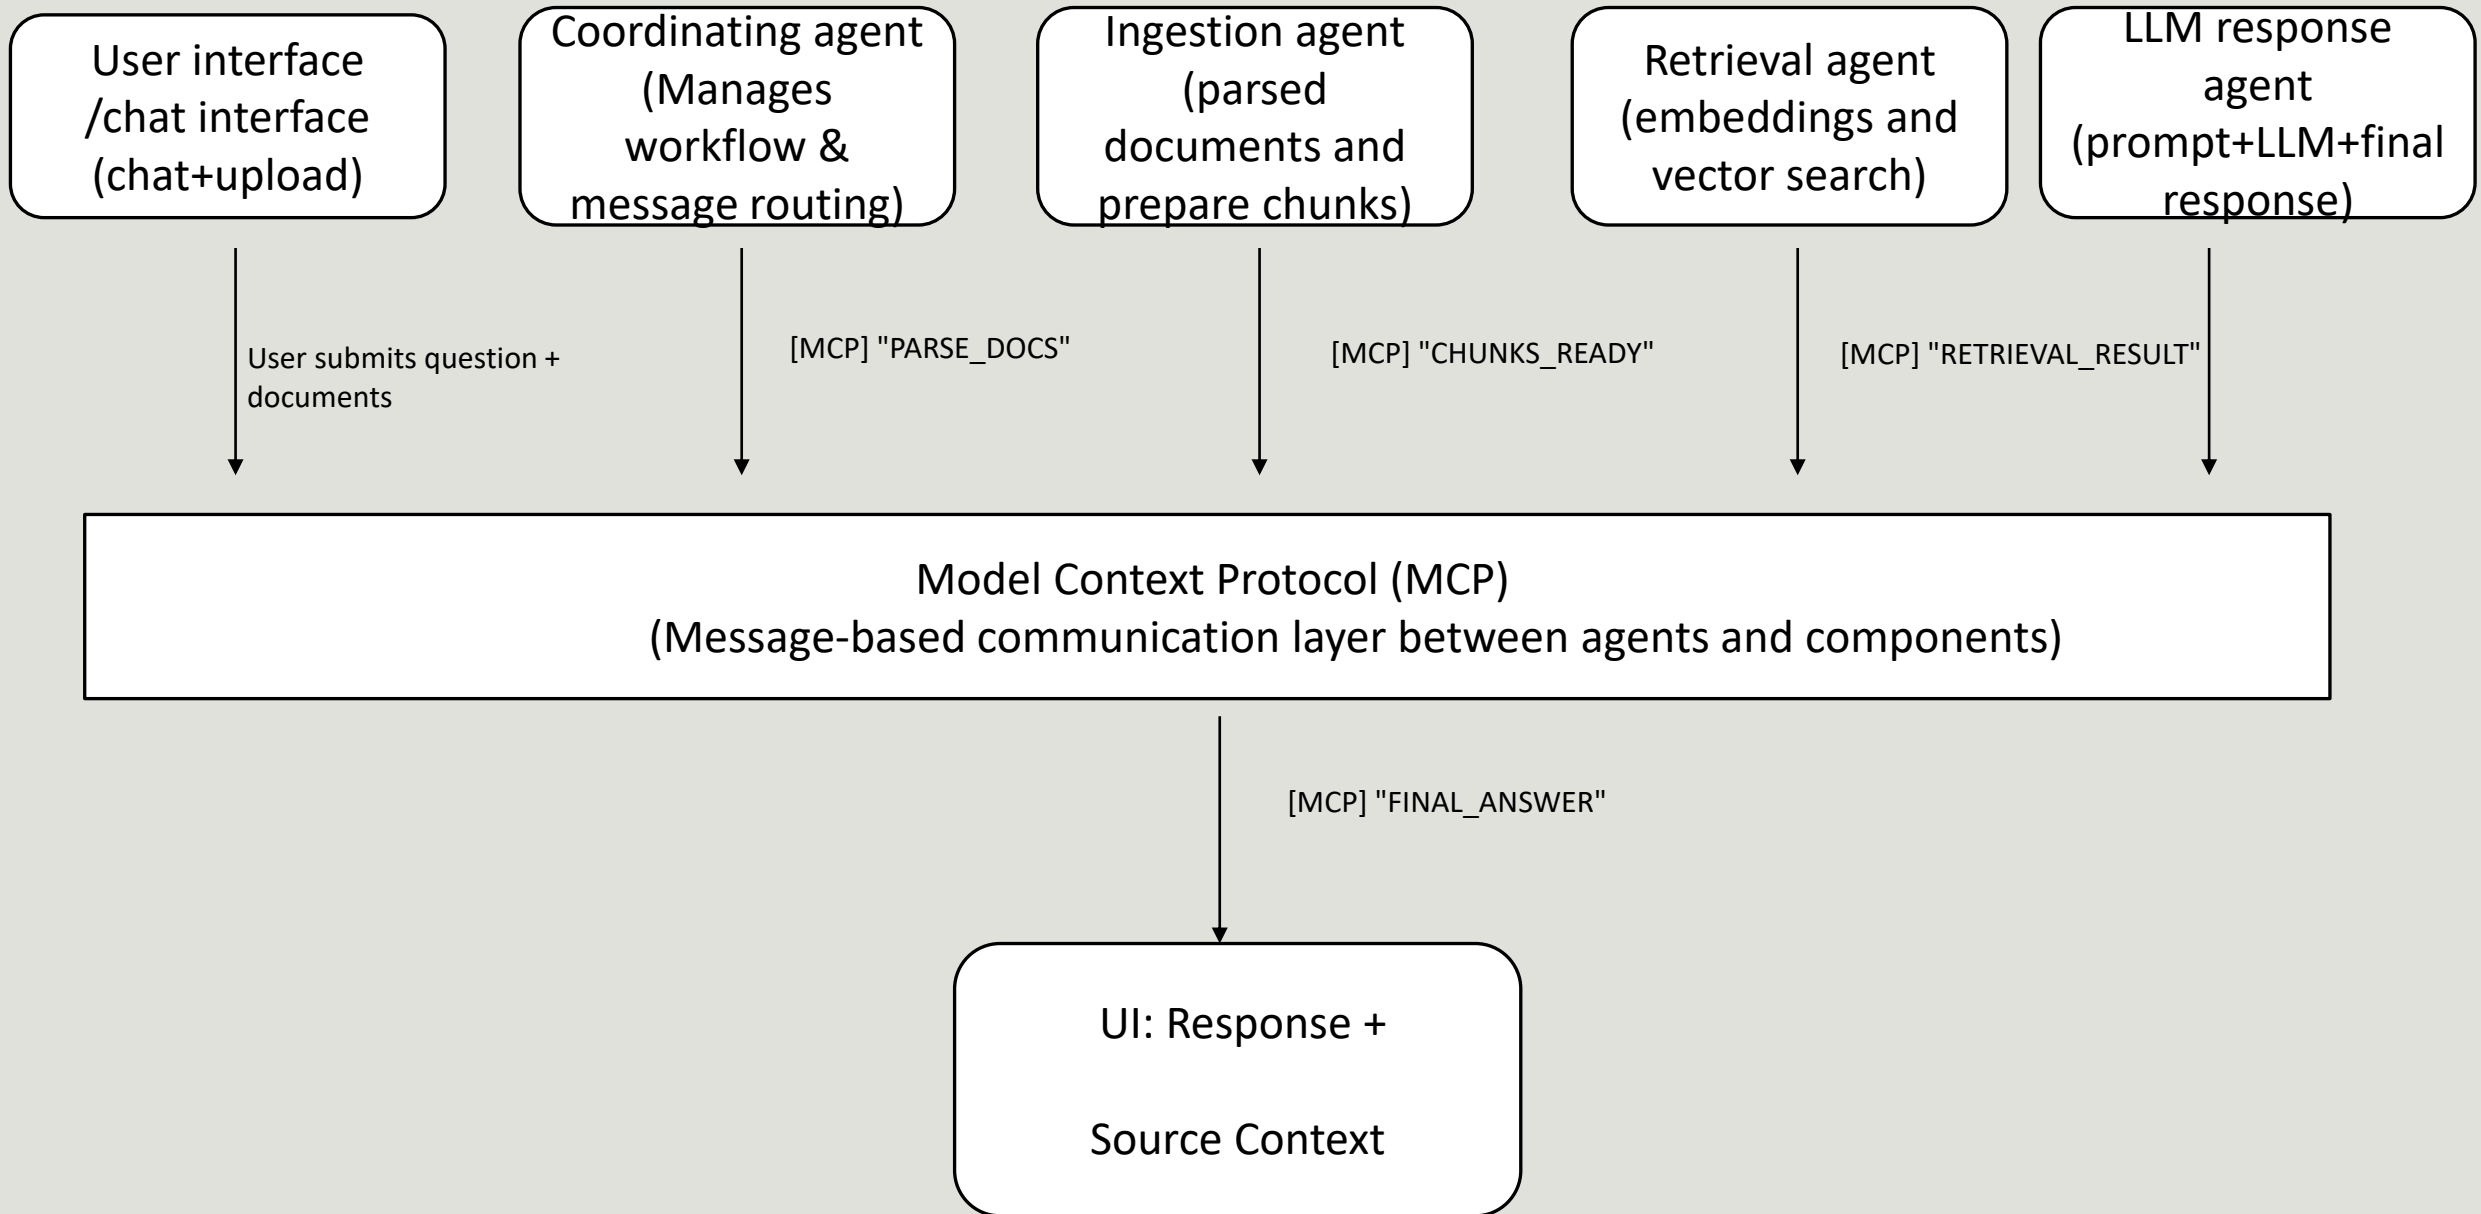


Agentic RAG Chatbot for Multi-Format Document QA

Introduction

This project presents the **Agentic RAG Chatbot**, an intelligent question-answering system designed to efficiently handle multi-format documents by leveraging the **Model Context Protocol (MCP)**. By combining retrieval-augmented generation with MCP, the chatbot dynamically accesses and processes diverse document types—including PDFs, Word files, spreadsheets, and more—to deliver accurate, context-aware responses. This innovative integration enables the system to maintain relevant context across user interactions and improve answer precision.



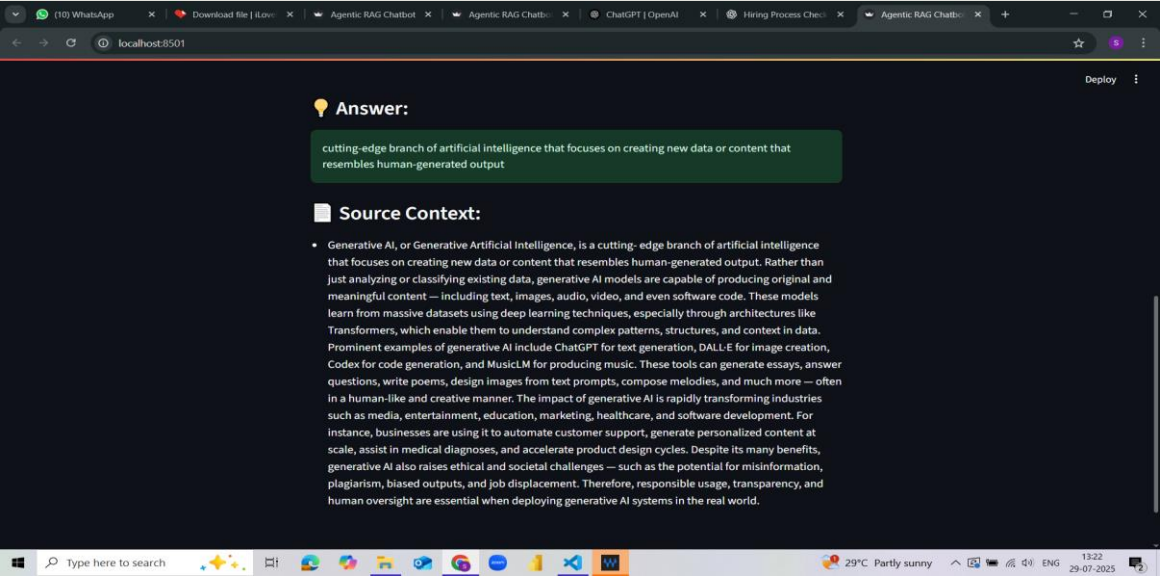
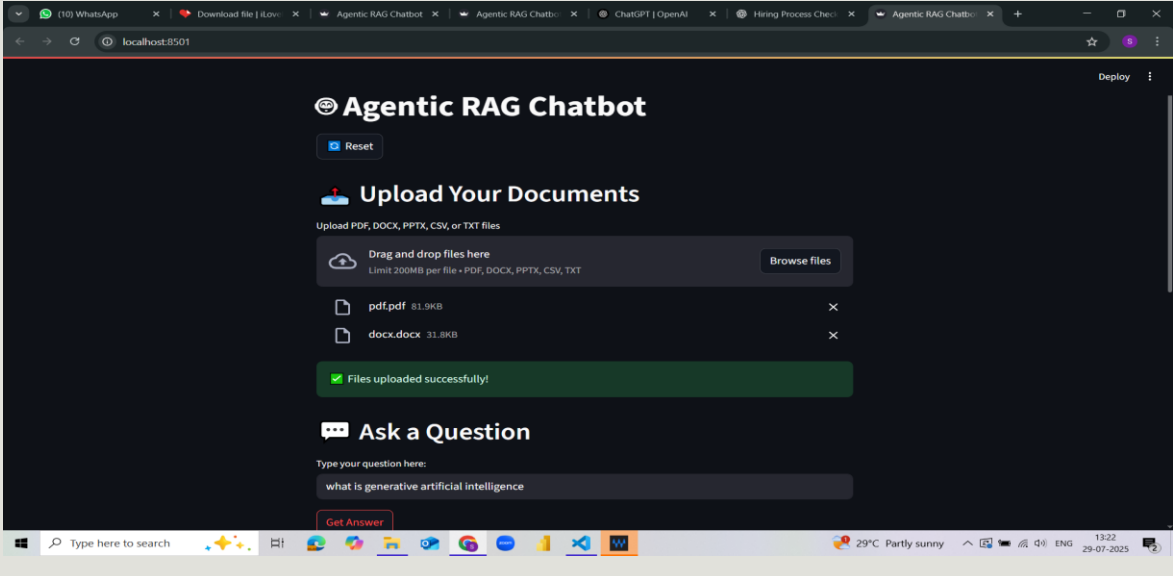
MCP protocol

```
import uuid
```

```
def create_mcp_message(sender, receiver, msg_type, payload, trace_id=None):  
    return {  
        "sender": sender,  
        "receiver": receiver,  
        "type": msg_type,  
        "trace_id": trace_id or str(uuid.uuid4()),  
        "payload": payload  
    }
```

Category	Technology	Purpose
Frontend (UI)	Streamlit	Chat interface, file uploads, user interaction
Backend Logic	Python	Core development language
Agents	Custom Python Classes	IngestionAgent, RetrievalAgent, LLMResponseAgent
Messaging Layer	Model Context Protocol (MCP)	Message-passing between agents (in-memory / REST)
Document Parsing	PyMuPDF, python-docx, python-pptx, pandas	Parsing PDF, DOCX, PPTX, CSV, TXT files
Embeddings	Hugging Face Transformers	Generating vector representations of chunks
Vector Store	Vector DB with FAISS	Storing and retrieving document embeddings
LLM	google/flan-t5-base	Answer generation using retrieved context
File Handling	Python OS / tempfile / base64	Upload, read, and manage user documents
Deployment	Streamlit	App deployment and demo hosting
Version Control	Git + GitHub	Code management and collaboration

Results:



Challenges Faced

Repetitive Content in Source Documents

Several documents contained repeated paragraphs or duplicated content across different formats (e.g., same text in PDF and PPTX).

This led to redundant chunks being passed to the LLM, increasing token usage and sometimes causing repetitive or verbose responses.

Solution: Implemented a deduplication step during document parsing to filter out repeated text before embedding and retrieval.